

## References

- Triplet Loss: [FaceNet: A Unified Embedding for Face Recognition and Clustering \(<https://arxiv.org/pdf/1503.03832.pdf>\)](https://arxiv.org/pdf/1503.03832.pdf)
- Sentence BERT: [entence-BERT: Sentence Embeddings using Siamese BERT-Networks \(<https://arxiv.org/pdf/1908.10084.pdf>\)](https://arxiv.org/pdf/1908.10084.pdf)

## Embeddings

From our module on [Transfer Learning \(Transfer\\_Learning.ipynb\)](#) and [CLIP \(CLIP.ipynb\)](#).

- we understand that a Sequential Deep Learning model
- has multiple layers prior to the task-specific "head" (e.g., Classifier)
  - that create alternate representations of the Input
  - of increasing complexity as we go Deeper (closer to the Head)

The purpose of the layers from the Input to the Head

- is to create a representation
- from which the Head can correctly solve the task (e.g., predict the label)

These alternate representations are called *embeddings* of the Input.

We also hypothesize that

- embeddings that are close to the input
  - represent "syntax" concepts embeddings that are close to the Head
    - represent "semantic" concepts
    - but if we are too close to the Head
      - embeddings are over-specialized to the Source task used for training

We wish to understand behavior of embeddings visually.

Given a number of inputs

- we can compute the embedding of each
- and plot them in embedding space
- in order to understand the structure of embedding space

Here is a plot of the embeddings

- of a subset of the 10 digits

Although embeddings are typically high dimension

- we plot in 2D space
  - define by the first 2 Principal Components
- as a practical matter

And here is the clustering of text articles across different classes.

</table>

Attribution: <https://joeddav.github.io/blog/2020/05/29/ZSL.html>  
[\(https://joeddav.github.io/blog/2020/05/29/ZSL.html\)](https://joeddav.github.io/blog/2020/05/29/ZSL.html)

It would seem

- that examples of the *same class* (e.g., digit or article topic)
- have embeddings that are close to one another
  - forming clusters in embedding space

This may explain why

- feeding embeddings into a Classifier head
- results in successful Classification



# Using Embeddings to solve a task without finetuning

The traditional way that we have used Embeddings is via Transfer Learning:

- keep a prefix of layers of a model trained for a Source task
  - use the embeddings created by the prefix
- append a new Head for a Target task
- train the Head on a small number of examples of the Target task
  - possibly also fine-tuning the weights in the prefix



## Zero-shot classification

A simple approach to Classification via embeddings is like K-nearest neighbor  
But, we could use the embeddings directly

- map an input  $\text{\textbackslash}x$  to its embedding  $\text{embed}(\text{\textbackslash}x)$
- starting with a dataset of examples for a Target task
- predict  $\hat{y}$ , the label of  $\text{\textbackslash}x$ 
  - map each input to its embedding

In this approach:

- as the label  $\text{\textbackslash}y^{\text{ip}}$  of an example  $\langle \text{\textbackslash}x^{\text{ip}}, \text{\textbackslash}y^{\text{ip}} \rangle$
- in a pre-defined set of labeled examples

$$[\langle \text{\textbackslash}x^{\text{ip}}, \text{\textbackslash}y^{\text{ip}} \rangle \mid 1 \leq i \leq m]$$

- we don't need to modify (or need access to the source code) the prefix layers whose embedding  $\text{embed}(\text{\textbackslash}x^{\text{ip}})$
- just treat the Source model as producing embeddings
  - is closest to  $\text{embed}(\text{\textbackslash}x)$

$$\text{embed}(\text{\textbackslash}x) \approx \text{embed}(\text{\textbackslash}x^{\text{ip}})$$

- so predict  $\hat{y} = \text{\textbackslash}y^{\text{ip}}$

This is motivated by the possibility

- that a new Target task
- can be solved *without training*
- just by comparing embeddings

# Semantic search

Want to create your own search engine ?

- create text embeddings for each document in a collection
- Note that this approach  
• create an embedding of your query

The document embeddings are computed once and stored. Computing a query embedding and comparing it to the document embeddings will give the correct result.

## Note

This is the basis for *Vector Stores*

- augmenting a LLM with your own data (e.g., GPT)

A simple example: facial (or image) recognition

- compare the embedding of an image
- with the embeddings of the fixed number of images for each class (e.g.

# Creating embeddings for similarity

Using the similarity of embeddings to solve a new Target task

- assumes that
- the embedding or "related" inputs
- are close to one another

We may get lucky and this will be the case for many tasks.

But we may need to train our embeddings

- with this as the specific goal
- by adding this as an objective of the Loss function



One such objective is the [Triplet Loss](https://arxiv.org/pdf/1503.03832.pdf) (<https://arxiv.org/pdf/1503.03832.pdf>)

Consider an input  $a$  (the "anchor")

## Example: Sentence Embeddings

• with related input  $p$  ("positive")

• with unrelated input  $n$  ("negative")

To illustrate, we show [Sentence BERT](https://arxiv.org/pdf/1908.10084.pdf) (<https://arxiv.org/pdf/1908.10084.pdf>)

Let

- fine-tunes the embeddings produced by BERT
- in order to make related sentences close in embedding space
- $s_a, s_p, s_n$  be the embedding produced by some layer, given input  $a$ ,
- $\|s - s'\|$  be a measure of the distance (inverse of similarity, always  $> 0$ ) between two embeddings  $s, s'$

The Triplet Loss objective is to *minimize*

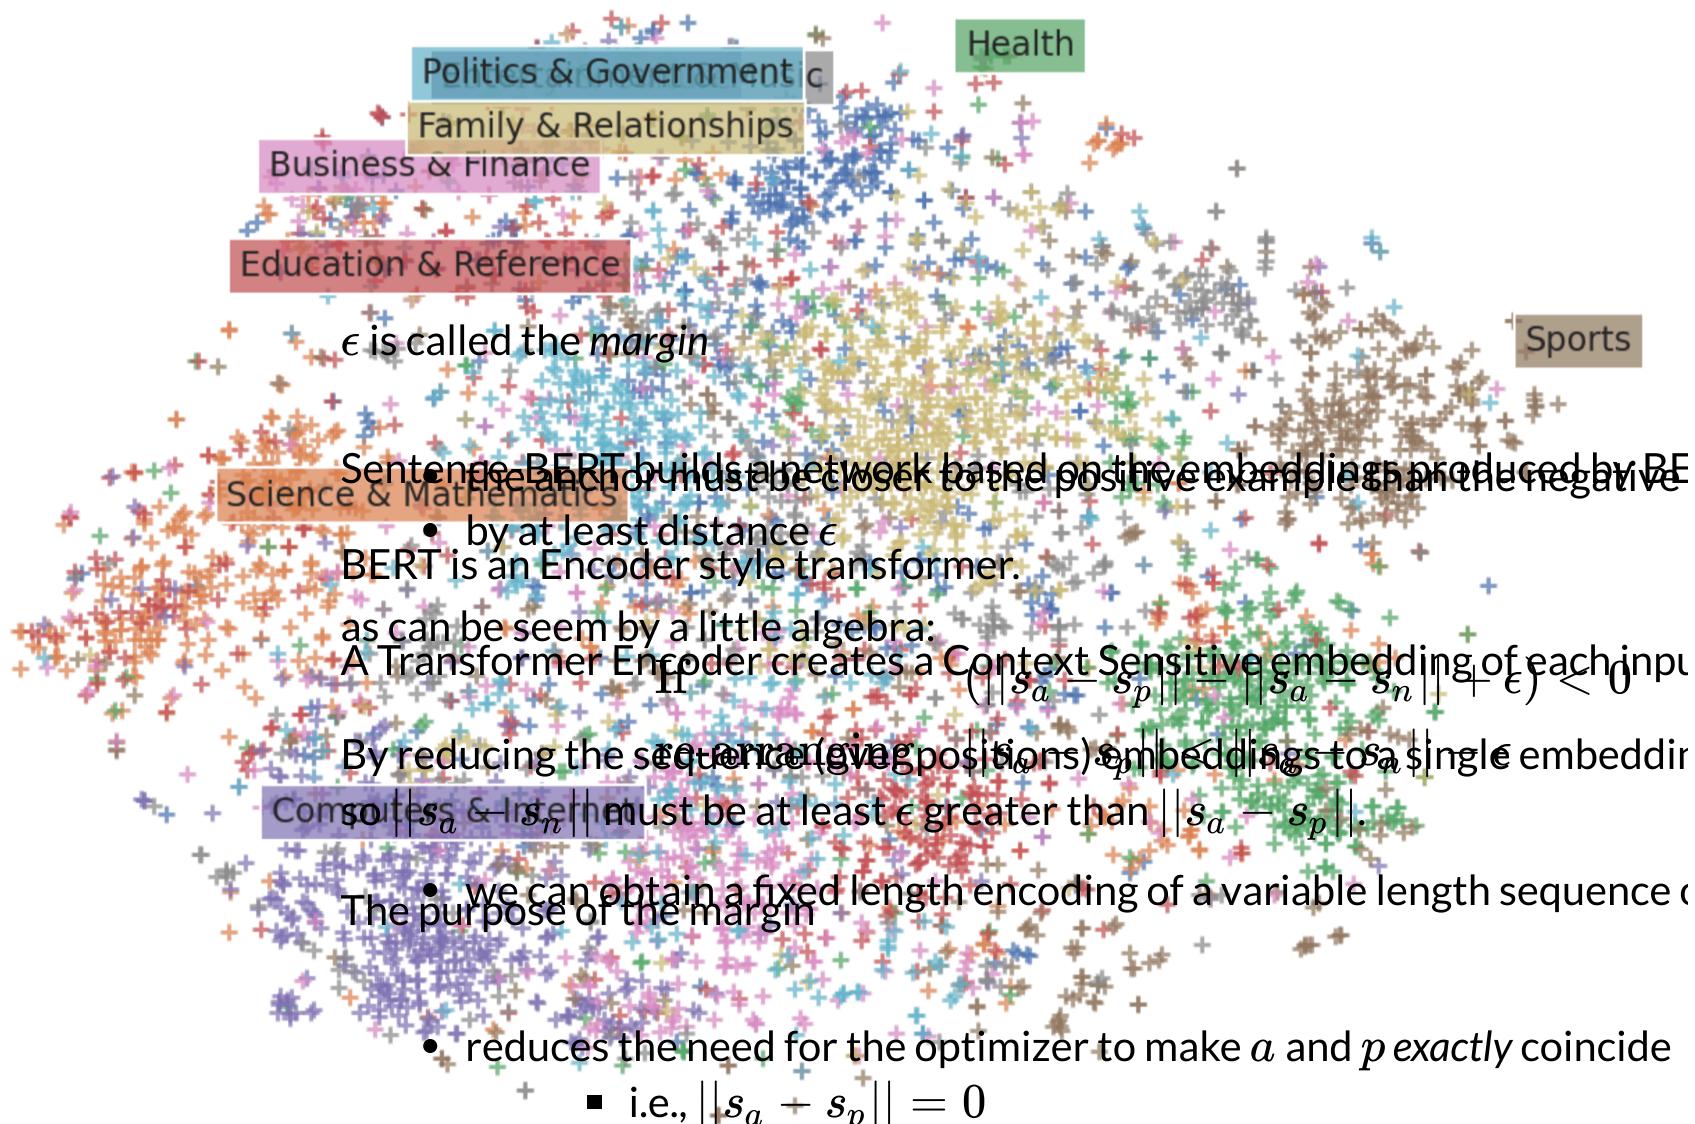
$$\max (||s_a - s_p|| - ||s_a - s_n|| + \epsilon, 0)$$

The loss is minimized when

- $s_a$  is close to  $s_p$
- $s_a$  is far from  $s_n$

That is the embedding for anchor

- $a$  is very similar to that for  $p$
- $a$  is very dissimilar to that for  $n$



&lt;/table&gt;

Attribution: [\(https://arxiv.org/pdf/1908.10084.pdf#page=3\)](https://arxiv.org/pdf/1908.10084.pdf#page=3)

The pre-trained BERT model is *shared* across two inputs: Sentence A and Sentence B.

- "weights are tied"

**Aside** BERT's weights are fine-tuned via the Triplet Loss objective

The sequence output of BERT is reduced by pooling (in this case),

Training a source task with the Triplet Loss to perform the reduction of a sequence to a single value

- Sentence A is embedded as  $u$
- Sentence B is embedded as  $v$
- Sentence C is embedded as  $w$

pooling (average over the embeddings) produces embeddings that are useful for tasks that compare similarity of embeddings

In the diagram on the right, the Triplet Objective uses a beginning/end "special" token (e.g.,  $\langle \text{CLS} \rangle$ ) to capture the entire sequence

- is recast as maximizing similarity (cosine distance)
- rather than minimizing distance

## Aside

The diagram on the left is for producing embeddings for a specific task

- entailment
  - Does Sentence B logically follow from Sentence A
- and hence is expressed as a Classification objective over labels  
 $\{”\text{Entail}”, ”\text{Does not entail}”\}$   
Here is the architecture

The inputs to the classifier are the concatenation of

- the embedding  $u$  of Sentence A
- the embedding  $v$  of Sentence B
- the difference in the embeddings

(Presumably these three inputs facilitate Classification)

The model is trained via batches that contain a mixture of

- Positive examples: Sentence A and Sentence B *are related* (anchor  $a$  and positive  $p$ )
- Negative examples: Sentence A and Sentence B *are un-related* (anchor  $a$  and positive  $n$ )

Triplet loss is minimized (or Utility maximized) in each batch.

# Performance

Here (<https://github.com/UKPLab/sentence-transformers/blob/master/docs/models/sts-models.md#performance-comparison>) is a comparison of Sentence Embedding with other methods

The Sentence Embedding (Universal Sentence Encoder) scores highest

- outperforms Word Embeddings (the two GloVe entries)
- it *greatly outperforms* the simple reduction methods used on plain BERT:
  - pooling (BERT as a service avg embeddings)
  - special <CLS> token (BERT as a service CLS vector)

## Note

The "sophisticated" BERT, when using simple reduction methods

- underperforms the "old school" word embeddings !



In [6]:

Done