

## References

- Rethinking the Role of Demonstrations: What Makes In-Context Learning Work? (<https://arxiv.org/pdf/2202.12837.pdf>)

# What makes In-context Learning work ?

- [blog \(<http://ai.stanford.edu/blog/understanding-incontext/>\)](http://ai.stanford.edu/blog/understanding-incontext/)
- [paper \(<https://arxiv.org/pdf/2202.12837.pdf>\)](https://arxiv.org/pdf/2202.12837.pdf)
  - more empirical
  - various models
    - MetalCL: trained with InContextLearning objective
    - 2 methods: Direct vs Channel ???
  - gold-label vs random (uniform sampling) label: ground-truth not necessary
    - gold improves over zero shot
    - random: small decrease vs gold
      - **very** small for MetalCL
    - sampling for true label distribution: smaller decrease

# How does In Context Learning work ?

*In-context Learning* describes a means of using a fixed LLM to solve a task

- by supplying some number  $k$  of *exemplars* (or *demonstrations*) of the new task
$$\langle \backslash \mathbf{x}^{(1)}, \backslash \mathbf{y}^{(1)} \rangle, \dots, \langle \backslash \mathbf{x}^{(k)}, \backslash \mathbf{y}^{(k)} \rangle$$
- as a context  $C$ 
  - that describes the new task's relationship between input  $\backslash \mathbf{x}^{\text{ip}}$  and output  $\backslash \mathbf{y}^{\text{ip}*}$
- and presenting a prompt  $\backslash \mathbf{x}$  to the model
- expecting the model to produce a  $\hat{\backslash \mathbf{y}}$
- that is the correct "response" to the new task on input  $\backslash \mathbf{x}$

So the prompt (context plus example's feature  $\backslash \text{x}$ ) might look like

Input:  $\backslash \text{x}^{(1)}$

Output:  $\backslash \text{y}^{(1)}$

:

Input:  $\backslash \text{x}^{(k)}$

Output:  $\backslash \text{y}^{(k)}$

Input:  $\backslash \text{x}$

Output:

and expect the continuation to be the prediction  $\hat{\text{y}}$  corresponding to test input  $\backslash \text{x}$

For example:

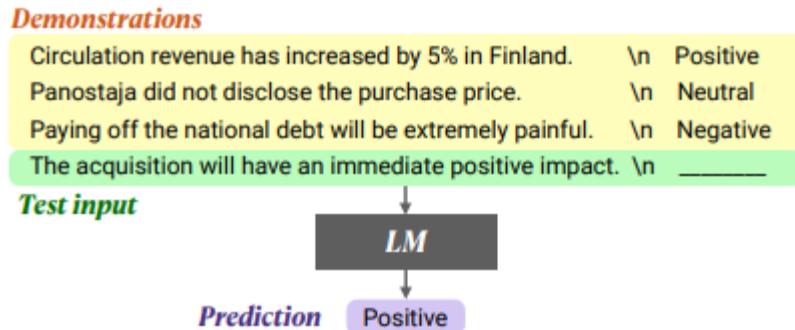


Figure 2: An overview of in-context learning. The demonstrations consist of  $k$  input-label pairs from the training data ( $k = 3$  in the figure).

Attribution: <https://arxiv.org/pdf/2202.12837.pdf#page=2>  
[\(https://arxiv.org/pdf/2202.12837.pdf#page=2\)](https://arxiv.org/pdf/2202.12837.pdf#page=2)

In-Context Learning uses a pre-trained LLM and the trick of using the Universal Text API

- to turn the new task
- into a text-continuation ("predict the next") task

It appears to be a way

- of extending a LM
- *without* further training
  - as opposed to Fine-Tuning
- since
  - the exemplars are given at *test* time
  - no parameter updates to the LLM occur

But why should this work ?

More interestingly

- what is a good theory
- and how can we test it

We will present a [paper \(<https://arxiv.org/pdf/2202.12837.pdf>\)](https://arxiv.org/pdf/2202.12837.pdf) that attempts to present some insights into the process.

# Testing some theories

In order to test a theory

- various aspects of the exemplars are proposed as variables
- one variable at a time is perturbed
- the effect of the perturbations is measured across a range of benchmarks
- **and compare to measurements before the perturbation**

The results are summarized in the following diagram

- that we will subsequently refer to for each experiment

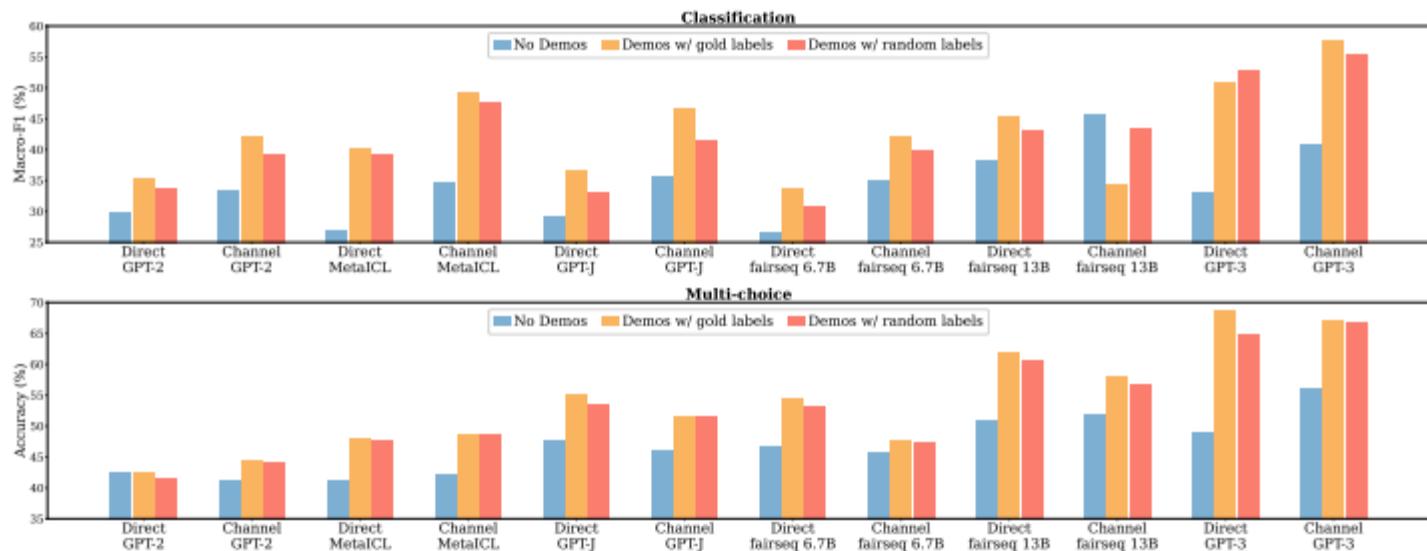


Figure 3: Results when using no-demonstrations, demonstrations with gold labels, and demonstrations with random labels in classification (top) and multi-choice tasks (bottom). The first eight models are evaluated on 16 classification and 10 multi-choice datasets, and the last four models are evaluated on 3 classification and 3 multi-choice datasets. See Figure 11 for numbers comparable across all models. **Model performance with random labels is very close to performance with gold labels** (more discussion in Section 4.1).

This chart shows the result of perturbations

- run across a variety of models
  - of sizes ranging from 774M to 175B parameters
- each experiment is averaged across multiple benchmarks

This summarizes a variety of experiments (to be described) involving the exemplars.

The 3 different colored bars represent variations on the exemplars

blue no exemplars

gold exemplars  $k = 16$

red exemplars with perturbed labels  $k = 16$

# **Zero shot verus $k \geq 1$ shot**

The first experiment measures the effect of the presence/absence of exemplars

- blue vs. orange

## **Conclusion**

$k \geq 1$  exemplars *improves performance* relative to zero-shot.

# Parts of the Context

The next set of experiments varies parts of the context (exemplars and prompt).

Given exemplars

$$\langle \backslash \mathbf{x}^{(1)}, \backslash \mathbf{y}^{(1)} \rangle, \dots, \langle \backslash \mathbf{x}^{(k)}, \backslash \mathbf{y}^{(k)} \rangle$$

the authors posit some salient characteristics

- the *input distribution*  $I$  from which the exemplar *features* are drawn  
 $\backslash \mathbf{x}^{(1)}, \dots, \backslash \mathbf{x}^{(k)}$
- the distribution  $L$  of the exemplar *labels*  $\backslash \mathbf{y}^{(1)}, \dots, \backslash \mathbf{y}^{(k)}$
- the feature/label mapping relationship  $M$ 
  - i.e., the pair of  $\backslash \mathbf{x}^{\text{ip}}$  and  $\backslash \mathbf{y}^{\text{ip}}$ , for  $1 \leq i \leq k$
- formatting
  - the encoding of the exemplars and prompt into  $\dot{\backslash \mathbf{x}}$

## Feature/label mapping relationship

- gold vs. red

Let  $\mathcal{C}$  denote the set from which exemplar labels are drawn.

In this experiment, replace

- correct label  $\textcolor{red}{\backslash y^{\text{ip}}}$  for exemplar  $i$
- with label  $\textcolor{red}{\tilde{y}^{\text{ip}}}$  drawn at random (uniformly) from  $\mathcal{C}$ .

That is, we preserve  $I$  and  $L$ , but break  $M$ .

## Conclusions

- Correct ("gold") labels improve performance over random labels
  - but not as much as expected, perhaps
- But Random labels *improves performance over no exemplars !*
  - "Ground truth" matters surprisingly little !

The fact that an *incorrect*  $M$  improves performance relative to no exemplars is surprising.

This suggests

- that the exemplars are used to infer the *task to be performed*
- once the task has been identified
  - the exemplar mis-labeling is ignored
- the model is able to perform the task as it was *trained* during training

See the [Signifier theory in the module](#)

[\(Prompt\\_Engineering\\_Suggestions.ipynb#Signifier:-direct-specification\)](#)

## Input distribution

This experiment measures a *shift in the distribution* of the exemplar features  $\mathbf{x}^{ip}$

In this experiment

- each exemplar input  $\mathbf{x}^{ip}$  is replaced by
- a random  $\mathbf{x}_{rand}^{ip}$  drawn from a text corpus *other than the one used for Training*

We note that this experiment *also* breaks the feature/label relationship  $M$

- we preserve the original labels  $\setminus \mathbf{y}^{\setminus ip}$  for exemplar  $i$
- which is not necessarily related to  $\setminus \mathbf{x}_{\text{rand}}^{\setminus ip}$

We can contrast the results of this experiment the effect of breaking  $M$  alone.

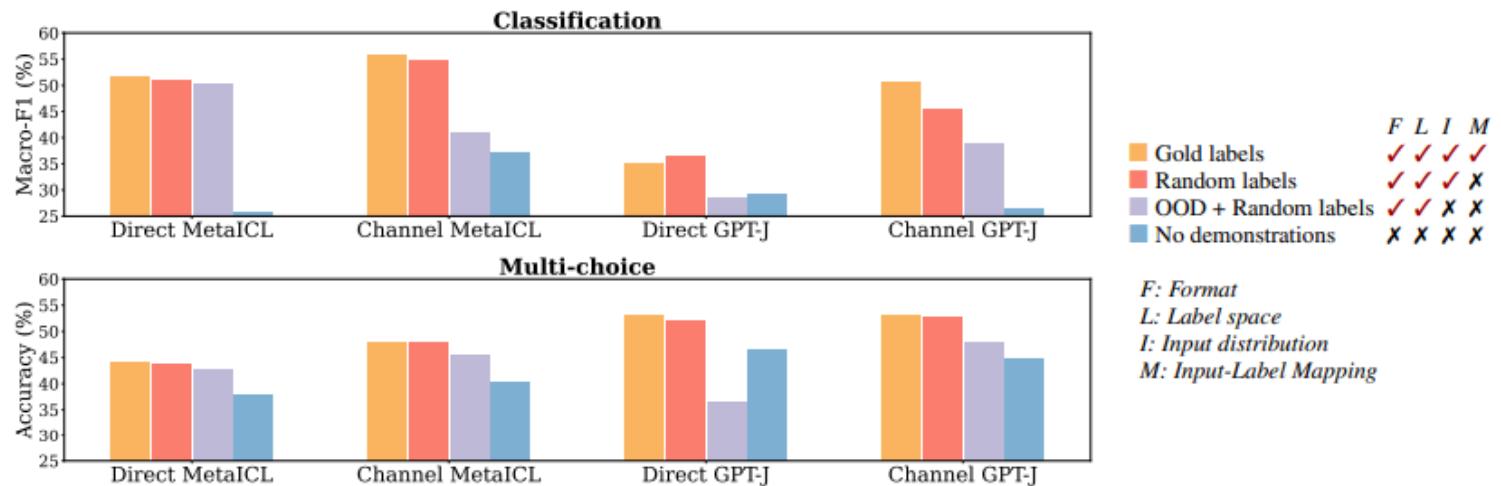


Figure 8: Impact of the distribution of the inputs. Evaluated in classification (top) and multi-choice (bottom). The impact of the distribution of the input text can be measured by comparing  $\text{Random labels}$  (red) and  $\text{OOD + Random labels}$  (purple). The gap is substantial, with an exception in Direct MetaICL (discussion in Section 5.1).

In the above diagram, compare

- the lavender (third bar from left): perturbed  $I$  and  $M$
- the red bar (second bar from left): perturbed  $M$  alone

## Conclusions

The  $M$  relationship is broken in both cases. But

- preserving the original distribution  $I$  of exemplar features
- improves performance relative to changing the distribution

Why might this be ?

The suggestion is that the model was trained with the LLM objective ("predict the next")

- from a training distribution
- and  $\textcolor{red}{x}_{\text{rand}}$  is from a *different* distribution
- so the model struggles on non-training input

# Output distribution

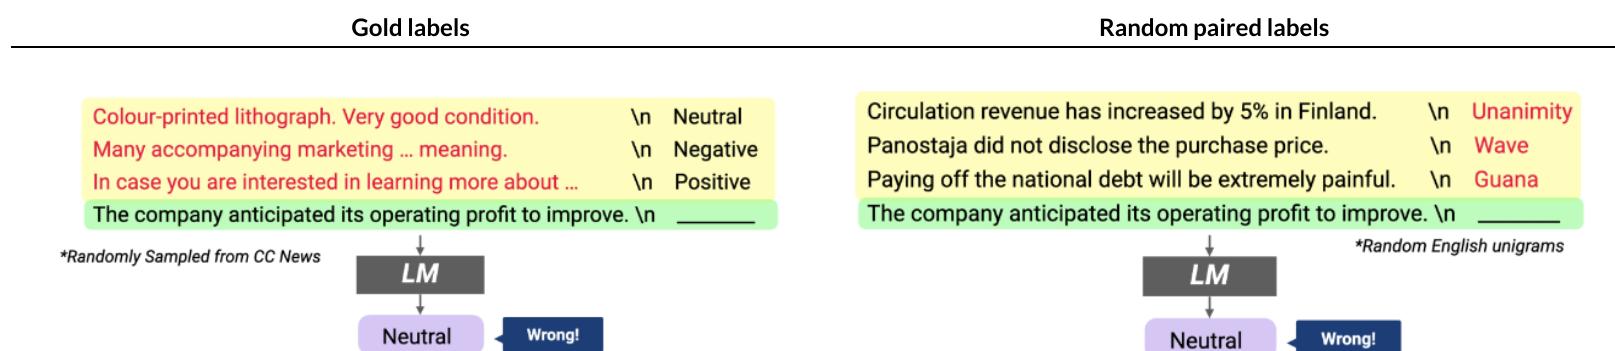
This experiment measures a *shift in the distribution* of the exemplar labels  $\backslash y^{\text{ip}}$

In this experiment

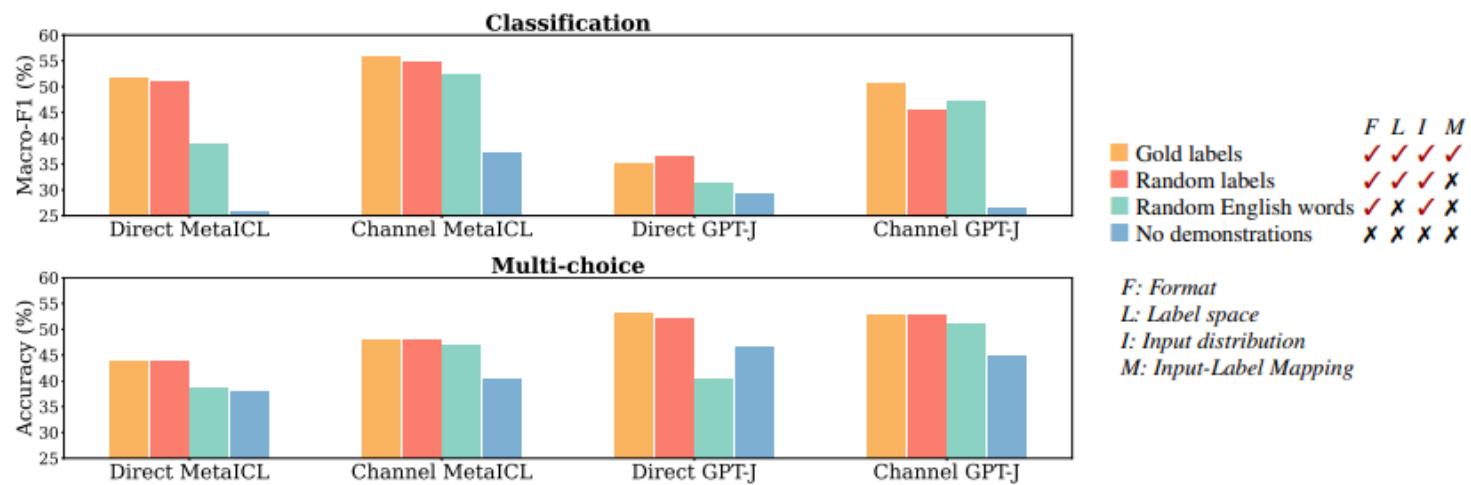
- each exemplar label  $\backslash y^{\text{ip}}$  (from  $\mathcal{C}$ ) is replaced by
- $\tilde{y}^{\text{ip}}$  a *random* English word from  $\mathcal{C}_{\text{rand}}$

Note that we *also break M*

- e.g.,  $\backslash y^{\text{ip}} = \backslash y^{(i')}$  does not imply  $\tilde{y}^{\text{ip}} = \tilde{y}^{(i)}$



Attribution: <http://ai.stanford.edu/blog/understanding-incontext/>  
[\(http://ai.stanford.edu/blog/understanding-incontext/\)](http://ai.stanford.edu/blog/understanding-incontext/)



In the above diagram, compare

- the red bar: random labels
- the turquoise bar (third from left): random words

The difference: Although we break  $M$  in both cases

- in the "random labels" case: the labels are chosen from the correct output distribution  $\mathcal{C}$
- in the "random words" case: the labels come from a distribution other than  $\mathcal{C}$

## Conclusions

The  $M$  relationship is broken in both cases. But

- preserving the original distribution  $L$  of exemplar labels
- improves performance relative to changing the distribution of labels

# Formatting

In this experiment

- *format* is defined as the *pairing* of a feature and label within an exemplar
- not necessarily a *correct pairing*: mapping  $M$  not necessarily correct

One experiment is run

- with *only* exemplar features (and no exemplar labels):  $\backslash \text{x}^{(1)}, \dots, \backslash \text{x}^{(k)}$
- natural comparison is with experiment of correct format
  - $\backslash \text{x}^{\text{ip}}$
  - paired with random English words (from  $\mathcal{C}_{\text{rand}}$ ) as labels

A second experiment is run

- with *only* exemplar labels (and no exemplar features):  $\backslash y^{(1)}, \dots, \backslash y^{(k)}$
- natural comparison is with experiment of correct format
  - a random  $\backslash x_{\text{rand}}^{\text{ip}}$  drawn from a text corpus
  - paired with  $\backslash y^{\text{ip}}$

Both comparison experiments

- preserve the format: feature/exemplar pairs
- without preserving  $M$
- or the distribution  $I$  in the first case, and  $L$  in the second case

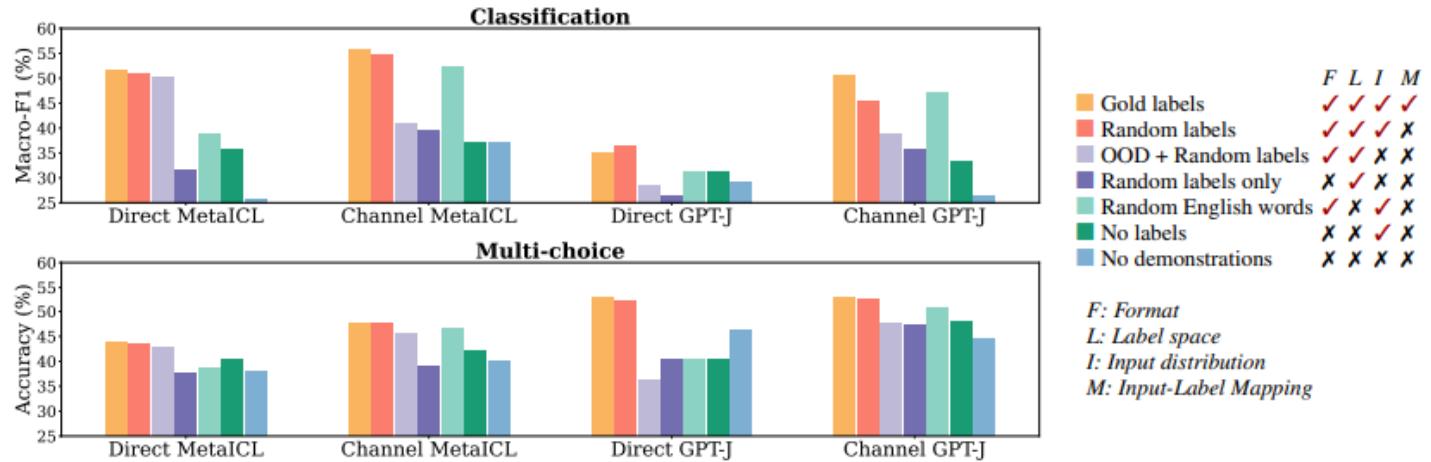


Figure 10: Impact of the format, i.e., the use of the input-label pairs. Evaluated in classification (top) and multi-choice (bottom). Variants of demonstrations without keeping the format (■ and ▨) are overall not better than no demonstrations (□). Keeping the format is especially significant when it is possible to achieve substantial gains with the label space but without the inputs (■ vs. ▨ in Direct MetaICL), or with the input distribution but without the labels (▨ vs. ▨ in Channel MetaICL and Channel GPT-J). More discussion in Section 5.3.

## Conclusions

- Not keeping the format has performance *on par* with **no demonstrations** at all
- Keeping the format retains most of the benefit achievable with either
  - correct  $I$  (but incorrect  $M$ )
  - or correct  $L$  (but incorrect  $M$ )

The suggestion is that correct format an important feature

- to enable the LLM to recognize the task from exemplars

## Exemplars that differ from the LLM

The on-line articles makes another interesting observation.

Observe the encoding of the exemplars

$$\langle \backslash \mathbf{x}^{(1)}, \backslash \mathbf{y}^{(1)} \rangle, \dots, \langle \backslash \mathbf{x}^{(k)}, \backslash \mathbf{y}^{(k)} \rangle$$

and example feature  $\backslash \mathbf{x}$

into the prompt

$$C, \backslash \mathbf{x}$$

The distribution of prompt  $C$ ,  $\backslash \text{x}$

- is probably *much different* than
- the distribution (Internet text documents) on which the LLM was trained

in several ways

- syntax
  - structured list of  $\backslash \text{x}^{\text{ip}}$ ,  $\backslash \text{y}^{\text{ip}}$  pairs
  - versus natural sentences
- coherence
  - consecutive exemplars may be from different *topics*
    - that illustrate the *concept* of the new task
  - consecutive natural sentences may share the same topic

The article posits that

- these encoding anomalies
- are low-frequency noise
- that the LLM is able to ignore
- providing there is more "signal" in the exemplars

# A theory of In-Context Learning

A more [theoretical paper](https://arxiv.org/pdf/2111.02080.pdf) (<https://arxiv.org/pdf/2111.02080.pdf>) and accompanying [online article](http://ai.stanford.edu/blog/understanding-incontext/) (<http://ai.stanford.edu/blog/understanding-incontext/>).

- combine these experimental insights
- into a theory
- and mathematical model of the theory
- that is consistent with the experimental results

The authors posit

- during training, the LLM learns "concepts", for example
  - abstract ideas
    - question answering
    - sentiment
  - plans
    - how to solve a multi-step task: travel directions

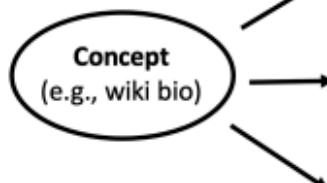
# Concepts

- 1. Pretraining documents**  
are conditioned on a  
**latent concept** (e.g.,  
biographical text)



Albert Einstein was a German theoretical physicist, widely acknowledged to be one of the greatest physicists of all time. Einstein is best known for developing the theory of relativity, but he also ....

- 2. Create independent examples from a shared concept.** If we focus on full names, wiki bios tend to relate them to nationalities.



| Input ( $x$ )       | Output ( $y$ ) | Delimiter                   |
|---------------------|----------------|-----------------------------|
| Albert Einstein was | German         | \n                          |
| Mahatma Gandhi was  | Indian         | \n                          |
| Marie Curie was     | ?              | ...brilliant?<br>...Polish? |

- 3. Concatenate examples into a prompt and predict next word(s). Language model (LM) implicitly infers the shared concept across examples despite the unnatural concatenation**

Albert Einstein was German \n Mahatma Gandhi was Indian \n Marie Curie was



Figure 1: In-context learning can emerge from modeling long-range coherence in the pretraining data. During pretraining, the language model (LM) implicitly learns to infer a latent **concept** (e.g., wiki bios, which typically transition between name (Albert Einstein) → nationality (German) → occupation (physicist) → ...) shared across sentences in a document. Although prompts are unnatural sequences that concatenate independent examples, in-context learning occurs if the LM can still infer the shared concept across examples to do the task (name → nationality, which is part of wiki bios).

Attribution: <https://arxiv.org/pdf/2111.02080.pdf#page=2>  
<https://arxiv.org/pdf/2111.02080.pdf#page=2>

The model's LLM "predict the next" training objective did not specify to goal of learning concepts.

But

- summarizing a large number of similar training documents (e.g., collection of biographies)
- in a parameters-constrained model
- logically suggests that concepts emerge as a way of reducing parameter usage

The authors suggest that the LLM's probability of outputting  $\text{\textbackslash}y$  given prompt  $\text{\textbackslash}x$  is formed by

$$\text{\textbackslash}pr\text{\textbackslash}y|\text{\textbackslash}x = \int_{c \in \text{Concepts}} \text{\textbackslash}pr\text{\textbackslash}y|\text{\textbackslash}x, c \text{\textbackslash}prc d(c)$$

That is, the output

- is the sum over all concepts
- of the probability of outputting  $\text{\textbackslash}y$  given prompt  $\text{\textbackslash}x$  and concept  $c$

Furthermore: the context (i.e., exemplars) of in-context learning

- helps the LLM identify the concept  $c$
- to which the prompt  $\backslash \text{x}$  implicitly refers

The experimental results seem to suggest that the exemplars

- don't need to be fully accurate
  - the model tolerates inaccurate mappings  $M$  between feature input space  $I$  and label space  $L$
- that correctly identifying  $I$  and  $L$  through the exemplars is
  - advantageous
  - but not completely necessary
- that the *format* of the exemplar
  - paired features and labels
  - is important

Under this theory

- the exemplars **are not teaching** new concepts
  - hence  $M$  can be inaccurate
- but serving to help the LLM **identify** a concept learned in training

That is, the encoded exemplars in  $C$

- are related to \prc

Once the concept  $c$  is identified, the output  $\hat{y}$  depends on

- the distributions  $I_{\text{train}}$  and  $L_{\text{train}}$ 
  - of features/labels
- the mapping  $M_{\text{train}}$ 
  - of features to labels

that were learned during **training**.

**NOT**

- on the  $I, L, M$  of the exemplars

That is

$$\hat{y} = \text{pr } y | \mathbf{x}, c \text{ concept } c$$

not

$$\hat{y} = \text{pr } y | C, \mathbf{x} \text{ exemplar context } C$$

In [2]: `print("Done")`

Done

