

Reasoning in Latent Space

Recall the behavior of LLM's that use reasoning/"Chain of Thought"

Given prompt

$$\backslash \textcolor{red}{x}_{(1:\bar{T})}$$

rather than *immediately* producing response

$$\backslash \textcolor{red}{y}_{(1:\bar{T})}$$

giving computation trace

$$\backslash \textcolor{red}{x}, \backslash \textcolor{red}{y}$$

- (dropping the sequence subscript that indexes tokens)

an LLM is trained to think "Step by Step"

- creating a sequence of steps
- the Chain of Thought ("reasoning" trace) $\backslash \textcolor{red}{rat}$
- enumerating sequential steps of a process that produces the response
$$\backslash \textcolor{red}{rat} = [\backslash \textcolor{red}{rat}_{(1)}, \dots, \backslash \textcolor{red}{rat}_{(\text{num_thoughts})}]$$
- where each $\backslash \textcolor{red}{rat}_{\backslash \text{tp}}$ is a thought represented as multi-token sequence

resulting in trace

\x, \rat, \v

All the tokens in sequences `\x`, `\rat`, `\y` come from the vocabulary `\Vocab`.

Thus, all the token sequences are in "Natural Language".

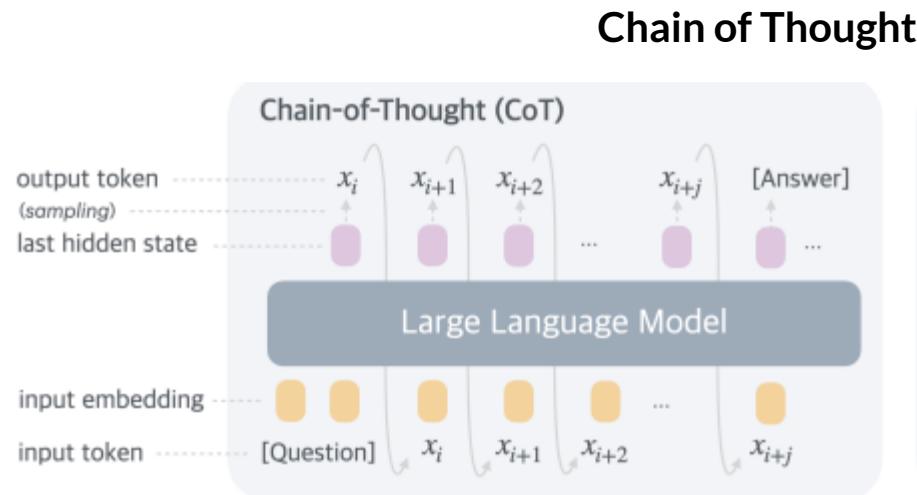
But the reasoning trace `\rat` is often hidden from the user.

- so `\rat` is not visible to the user

Does the LLM have to "think" in Natural Language ?

Here is an illustration of Chain of Thought reasoning using

- reasoning tokens
- limited to only Natural Language tokens



Attribution: <https://arxiv.org/pdf/2412.06769.pdf#page=2>

(Note: paper uses $\backslash x$ as a denotation of the output sequence vs. our standard of $\backslash y$)

Note

We are not quite precise in referring to the elements of the sequence as "tokens"

Technically:

- the elements of the depicted reasoning trace are the dense *embedding vectors* associated with the tokens
 - token: syntax; embedding: vector
 - one-to-one correspondence
- key point: elements of the reasoning are selected from a *finite vocabulary*
 - tokens/embeddings

Suppose we expanded the vocabulary of the reasoning trace

- to include embeddings
- that are **not restricted** to discrete categorical values from a finite Vocabulary
- but which are continuous vectors

Being continuous rather than discrete

- enables a richer space of "thoughts"
- which can't be mapped back to a token from a fixed vocabulary

This is called *Chain of Continuous Thought (Coconut)*.

What do these non-verbal continuous "tokens" represent ?

Some theories

- pictures/diagrams that "visualize" the answer
- super-position of two (or more) continuations
 - e.g., narrow down multiple choice from 4 to 2 possibilities -- keep those 2 possibilities alive

Chain of Thought

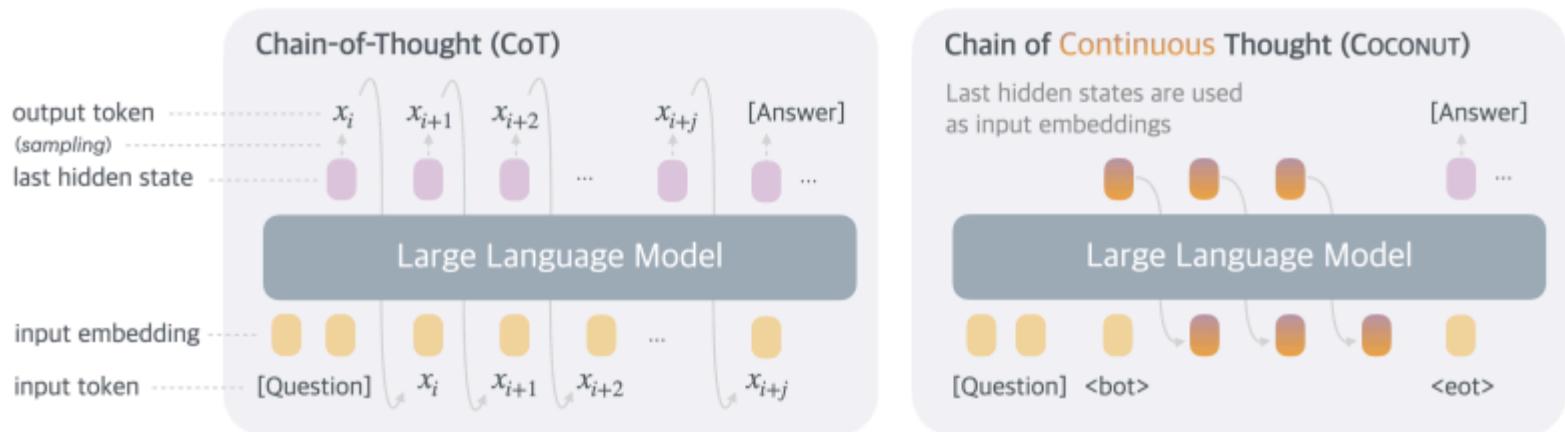


Figure 1 A comparison of Chain of Continuous Thought (CoCONUT) with Chain-of-Thought (CoT). In CoT, the model generates the reasoning process as a word token sequence (e.g., $[x_i, x_{i+1}, \dots, x_{i+j}]$ in the figure). CoCONUT regards the last hidden state as a representation of the reasoning state (termed “continuous thought”), and directly uses it as the next input embedding. This allows the LLM to reason in an unrestricted latent space instead of a language space.

Attribution: <https://arxiv.org/pdf/2412.06769#page=2>

Continuous thoughts

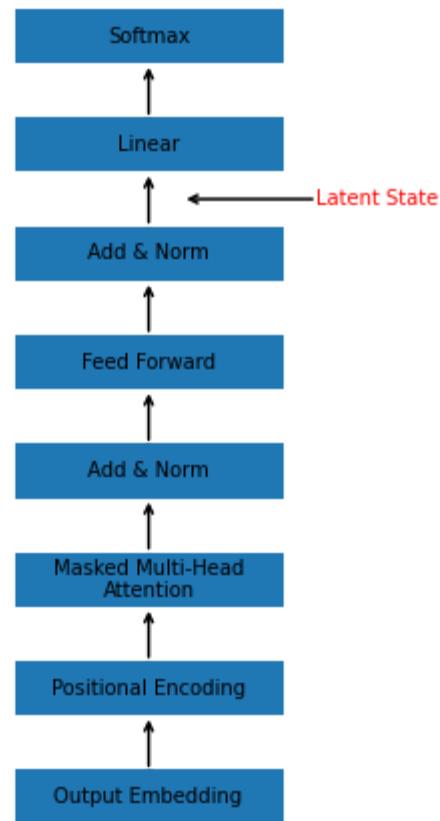
Where do the tokens that form the continuous thoughts come from ?

It is a simple adaptation of the Transformer architecture.

- see diagram below

In [5]: fig_transf

Out[5]:



The Transformer operates in Auto-regressive mode:

- outputs one token at a time
- at step t it's output \hat{y}_{tp}
- is appended to the previous output sequence $\hat{y}_{(1:t-1)}$
- which becomes the input for step $t+1$

The tokens (vectors of length $|\mathcal{V}|$, where \mathcal{V} is the Vocabulary of tokens) are

- *embedded* into vectors of length $d = d_{\text{model}}$
- and processed by the "body" of the Transformer
 - everything before the Linear layer

The output before the linear layer (referred to as the "Latent state")

- is a vector of length d
- which is *un-embedded* by the Linear layer into a vector of length $|\mathcal{V}|$
 - a probability vector over the finite, discrete set of language tokens \mathcal{V}
 - so a single token in \mathcal{V} can be output

We represent this schematic description of the Transformer in the left plot "Language Mode".

See diagram below.

The Embed block

- transforms discrete tokens into vectors whose length d is the path width of the Transformer

The Un-embed block

- transforms the latent state vector of length d back into a Language token
 - which must then be embedded before it can be used in the next time step

The Continuous Thought tokens

- are the final vectors of the Body
 - length d
 - *Latent state*

They are *not* un-embedded.

They can be fed *directly* into the Body at the next time step

- by-passing the Un-embed of step and the subsequent Embed of step (+1).

See the schematic diagram for "Latent Mode"

What does a Continuous token represent ?

Since a continuous token is not output to the user, it's vector representation does not have to be interpretable.

But the authors conducted an experiment

- decode the continuous token using the Un-embed
- resulting in a probability vector over the Language tokens in the NLP Vocabulary

Here is the input to the LLM

James decides to run 3 sprints 3 times a week.

He runs 60 meters each sprint.

How many total meters does he run a week ?

There are several paths that lead to the correct answer ($3 * 3 * 60 = 540$)

- First compute $(3 * 3)$, then multiple the result by 60
 - leading to outputting 9 as the first token of the reasoning trace
- First compute $(60 * 3)$
 - leading to outputting 180 as the first token of the reasoning trace

When the continuous thought vector is un-embedded

- it shows a probability distribution over the likely next tokens: 9 and 180

The authors interpret the thought vector as a "super-position" of possible next tokens.

Decoding a Continuous Thought

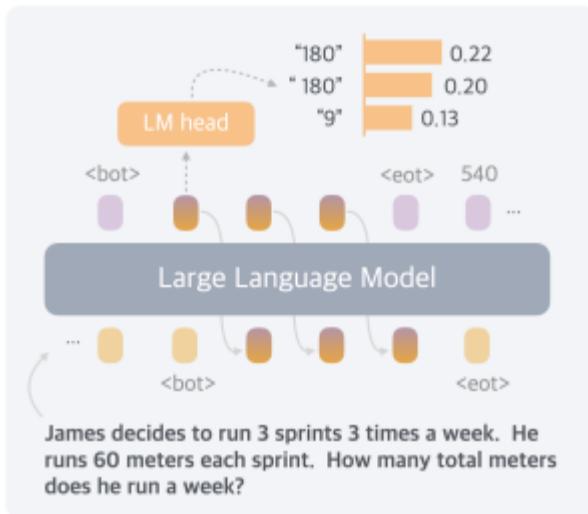


Figure 4 A case study where we decode the continuous thought into language tokens.

Attribution: <https://arxiv.org/pdf/2412.06769#page=7>

Absent "latent mode"

- the LLM would have to commit to **one** of the choices
 - e.g., "180" has the highest probability
 - in extending the reasoning trace
- by not committing immediately:
 - the LLM can "simultaneously" explore **both** possibilities
 - essentially allowing multiple traces (one beginning with "180", another beginning with "9")

The authors interpret this as the LLM

- performing a search across possible completions
- preserving possibilities
- until a later point in time

Training the LLM to reason in mixed language/latent mode

An LLM can be trained in CoT reasoning by using examples of reasoning traces

\x, \rat, \y

But this works only if the tokens of every thought in \rat are *language* tokens.

How do we train when \rat is a mixture of language and continuous tokens ?

- the continuous tokens are *defined* by training; not pre-specified

Clarification

Trace `\rat` consists of

- a number of steps
- each step consists of c "thoughts" (tokens: either language or continuous)

The authors use a clever technique:

Multi-stage training curriculum

- different training set at each stage
- at stage k
 - replace the first k steps ($k * c$ Natural Language tokens)
 - by continuous steps ($k * c$ continuous thoughts)
- stage 0 is traditional CoT training examples

At stage 0, the LLM learns how to reason.

- calculate a loss per language token in \rat
- guiding the learning-to-reason process

At subsequent stages, it learns to derive continuous tokens incrementally

- increasing k , the number of continuous thoughts

By proceeding incrementally

- the model does not "un-learn" traditional CoT in language tokens

The loss for continuous tokens

- is **masked** (not calculated)
 - no actual target with which to compare the prediction
- the loss of the language tokens following the continuous tokens
 - is reduced
 - when the preceding continuous tokens contribute to the prediction of the language tokens

The underlined elements in the diagram are the ones for which a loss is calculated.

Note

- each thought is a sequence of up to c tokens

Continuous tokens bracketed by `<bot>`, `<eot>`

Decoding a Continuous Thought

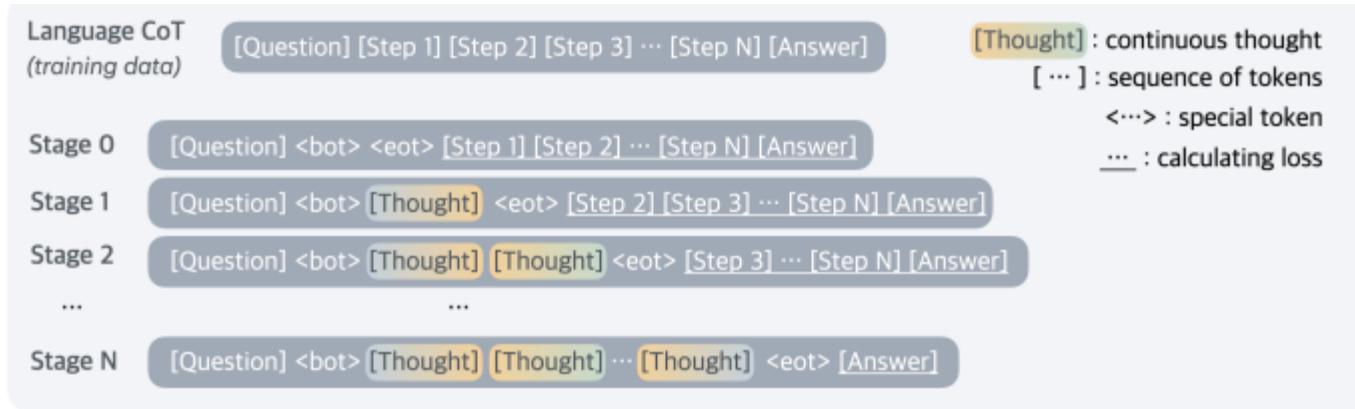


Figure 2 Training procedure of Chain of Continuous Thought (CoCoNUT). Given training data with language reasoning steps, at each training stage we integrate c additional continuous thoughts ($c = 1$ in this example), and remove one language reasoning step. The cross-entropy loss is then used on the remaining tokens after continuous thoughts.

Attribution: <https://arxiv.org/pdf/2412.06769.pdf#page=4>

Why is Latent Space potentially powerful

The reasoning paradigm of "think before you speak" is very powerful

- CoT leads to better outcomes
- by (potentially) hiding the thinking steps from the user
 - enables model to
 - reflect and revise
 - avoiding prematurely committing to a final answer

Allowing the thinking to occur in latent space

- superposition of multiple answers: no premature commitment to final answer
- non-verbal thoughts
 - images (e.g., a map enables one to think in relative position rather than precise directions)

Recall our original model for dealing with sequences

- Recurrent models with latent state
 - control of Finite State Machine
 - long term memory (LSTM)

Do the (hidden) "thoughts" of CoT have the same effect ?

- Hidden thought: a reminder of what to do next/later (control)

In [8]: `print("Done")`

Done

