

# Straight Through Estimator

## Motivation

The key to training a Neural Network

- is Gradient Descent
- implemented through Back Propagation

But there are conditions in which Back Propagation breaks down.

Let's illustrate the issue.

For notational consistency

- we will explain Back Propagation by equating operators and layers
- the operator for layer  $\ll$  maps input  $\textcolor{red}{y}_{(\ll-1)}$  to output  $\textcolor{red}{y}_{\text{lp}}$
- this works for operators organized in non-layered architecture as well

Our eventual goal is to compute the gradient of the Loss with respect to the weights  
 $\mathbf{W}_{\text{llp}}$  of each layer  $\ll$

$$\frac{\partial \text{loss}}{\partial \mathbf{W}_{\text{llp}}}$$

Back Propagation achieves this

- through repeated use of the Chain Rule
- starting from the deepest layer (head) and proceeding to the shallowest layer (input)

We define the Loss Gradient for layer  $i$  as

$$\text{loss}'_{\text{llp}} = \frac{\partial \text{loss}}{\partial \mathbf{y}_{\text{llp}}}$$

It is the derivative of  $\text{loss}$  with respect to the output of layer  $\ll$ , i.e.,  $\mathbf{y}_{\text{llp}}$ .

The desired gradient to update the weights follows by the chain rule using

- the Loss Gradient  $\text{loss}'_{\text{llp}}$
- the local gradient  $\frac{\partial \mathbf{y}_{\text{llp}}}{\partial \mathbf{W}_{\text{llp}}}$

$$\frac{\partial \text{loss}}{\partial \mathbf{W}_{\text{llp}}} = \text{loss}'_{\text{llp}} * \frac{\partial \mathbf{y}_{\text{llp}}}{\partial \mathbf{W}_{\text{llp}}}$$

Back propagation inductively updates the Loss Gradient from the output of layer  $\ll$  to its inputs (e.g., prior layer's output  $\text{\textbackslash}y_{(\ll-1)}$ )

- Given  $\text{\textbackslash}loss'_{\text{llp}}$
- Compute  $\text{\textbackslash}loss'_{(\ll-1)}$
- Using the chain rule

$$\begin{aligned}\text{\textbackslash}loss'_{(\ll-1)} &= \frac{\partial \text{\textbackslash}loss}{\partial \text{\textbackslash}y_{(\ll-1)}} \\ &= \frac{\partial \text{\textbackslash}loss}{\partial \text{\textbackslash}y_{\text{llp}}} \frac{\partial \text{\textbackslash}y_{\text{llp}}}{\partial \text{\textbackslash}y_{(\ll-1)}} \\ &= \text{\textbackslash}loss'_{\text{llp}} \frac{\partial \text{\textbackslash}y_{\text{llp}}}{\partial \text{\textbackslash}y_{(\ll-1)}}\end{aligned}$$

The loss gradient "flows backward", from  $\text{\textbackslash}y_{(L+1)}$  to  $\text{\textbackslash}y_{(1)}$ .

This is referred to as the *backward pass*.

That is:

- the upstream Loss Gradient  $\backslash \text{loss}'_{\text{llp}}$
- is modulated by the local gradient  $\frac{\partial \backslash \mathbf{y}_{\text{llp}}}{\partial \backslash \mathbf{y}_{(\ll -1)}}$
- where the "layer" is the operation transforming input  $\backslash \mathbf{y}_{(\ll -1)}$  to output  $\backslash \mathbf{y}_{\text{llp}}$

The problematic issue for Back-Propagation is the local gradient term

$$\frac{\partial \text{\\}y_{\text{\\}llp}}{\partial \text{\\}y_{(\ll -1)}}$$

that relates the output of an operation (layer) to its input.

What happens when the operation

- is non-differentiable
- or has zero derivative almost everywhere
- is non-deterministic (e.g., `tf.argmax` when two inputs are identical)

Either

- the backward gradient flow is killed (zero derivative)
- or can't be computed (non-differentiable or non-deterministic) **analytically**

## Solution: Straight Through Estimator (STE)

Suppose for the purpose of Back Propagation only we replace the problematic

$$\frac{\partial \textcolor{red}{y}_{\text{llp}}}{\partial \textcolor{red}{y}_{(\ll -1)}}$$

with the gradient of a proxy function  $\tilde{y}$

$$\frac{\partial \tilde{y}_{\text{llp}}}{\partial \textcolor{red}{y}_{(\ll -1)}}$$

That is, for the purpose of computing gradients

- we treat the operator as mapping  $\textcolor{red}{y}_{(\ll -1)}$  to  $\tilde{y}_{\text{llp}}$  rather than  $y_{\text{llp}}$

A common choice for the proxy function is the *identity function*

- operator implements  $\tilde{\mathbf{y}}_{\text{llp}} = \mathbf{y}_{(\ll -1)}$
- hence the formerly problematic gradient

$$\frac{\partial \mathbf{y}_{\text{llp}}}{\partial \mathbf{y}_{(\ll -1)}}$$

is replaced by

$$\begin{aligned} \frac{\partial \tilde{\mathbf{y}}_{\text{llp}}}{\partial \mathbf{y}_{(\ll -1)}} &= \frac{\partial \mathbf{y}_{(\ll -1)}}{\partial \mathbf{y}_{(\ll -1)}} \quad \text{since proxy implements } \tilde{\mathbf{y}}_{\text{llp}} = \mathbf{y}_{(\ll -1)} \\ &= 1 \end{aligned}$$

The Loss Gradient flows through the operator unchanged !

Allowing the Gradient to flow backwards unchanged allows us to construct something called a *Straight Through Estimator*

- treat problematic layer  $\ll$  as a pass-through of input  $\textcolor{red}{y}_{(\ll-1)}$  to  $\textcolor{red}{y}_{\text{llp}}$

A consequence of using a Straight Through Estimator is that

- Back propagation does not compute the true gradient
$$\frac{\partial \textcolor{red}{\text{loss}}}{\partial \textcolor{red}{W}_{\text{llp}}}$$
- but results in a value  $g_i$  (referred to as the *coarse gradient*)
  - not clear whether  $g_i$  is the gradient of any true function

This seems disturbing at first.

But recall that the purpose of computing the true gradient

- is to update weights  $\mathbf{W}_{\text{llp}}$

$$\mathbf{W}_{\ll} = \mathbf{W}_{\text{llp}} - \alpha * \frac{\partial \text{loss}}{\partial \mathbf{W}_{\text{llp}}}$$

- As long as the direction  $g_i$  has a non-zero correlation with  $\frac{\partial \text{loss}}{\partial \mathbf{W}_{\text{llp}}}$ 
  - the weight update step will move in the direction of reducing the Loss

Hence, the Straight Through Estimator can be used for the purpose of Gradient Descent.

# Proxy functions

Under [certain assumptions \(<https://arxiv.org/pdf/1903.05662.pdf#page=5>\)](https://arxiv.org/pdf/1903.05662.pdf#page=5).

- a non-zero correlation can be established between the true and coarse gradients
- for some specific proxy functions

Interestingly enough: the Identity function is **not** [one of those functions \(<https://arxiv.org/pdf/1903.05662.pdf#page=8>\)](https://arxiv.org/pdf/1903.05662.pdf#page=8) !

- ReLU and Clipped ReLU *are* such functions
  - these are more commonly used as Activation Functions
  - but here they just map an input to an output

Nonetheless: the Identity function is commonly used as the proxy

- it may initially lead to decreased Loss
- but will stop decreasing the loss as it approaches a local minimum

# Implementing a Straight Through Estimator in TensorFlow

## Stop Gradient operator

The *Stop Gradient* operator `sg` in TensorFlow

- acts as the identity operator on the Forward Pass

$$\text{sg}(\backslash \mathbf{x}) = \backslash \mathbf{x}$$

- But on the Backward Pass of Backpropagation: *it stops the gradient from flowing backwards*

$$\frac{\partial \text{sg}(\backslash \mathbf{x})}{\partial \backslash \mathbf{y}} = 0 \text{ for all } \backslash \mathbf{y}$$

We can use the Stop Gradient operator

- to prevent the computation of a gradient for a problematic operation/layer
- but we must take an extra step to allow the Loss Gradient to flow backwards through the problematic operation/layer

# Implementing Straight Through Estimation using Stop Gradient

Consider the implementation of an operator ProblemOp

- taking input `in`
- and producing output `out`
- defined by

```
class ProblemOp(layers.Layer):
    def call(self, in):
        # Computation of out:
        #   Problem: RHS of definition of result is NOT differentiable
        result = ...

        # Straight-through estimator.
        out = in + tf.stop_gradient(result - in)

    return out
```

n.b., the `call` method is what implements the Forward Pass in TensorFlow.

The line of code

```
out = in + tf.stop_gradient(result - in)
```

is the implementation of the \*Straight Through Estimator\*\*.

It seems odd: mathematically, it just copies `result` to `out`

## The line of code

- connects the output of the operator (the LHS of the assignment)  
 $\text{\textcolor{red}{y}}_{\text{llp}} = \text{out}$
- to the input of the operator (appearing on the RHS of the assignment)  
 $\text{\textcolor{red}{y}}_{(\ll -1)} = \text{in}$
- causing the Loss Gradient to flow from `out` to `in` during Back Propagation
  - unchanged
  - because the `sg` operator creates a zero gradient during Back Propagation

$$\frac{\partial \text{tf.stop\_gradient}(\text{result} - \text{in})}{\partial \dots} = 0$$

Thus, a Straight Through Estimator using an Identity proxy function has been created by the odd-looking statement.

Note too that the line of code

- **also** connects the output (LHS of assignment) to **result** (appearing on RHS of assignment)
- but, because of Stop Gradient
  - no gradient flows from **out** to **result**

Putting the problematic operation (calculation of **result**) within a Stop Gradient

- solves the potential issue of computing a problematic gradient

More formally:

Let's interpret this line of code as a layer  $\ll$

$$\begin{aligned}\text{\color{red} \textbackslash y\_llp} &= \text{in} &+ \text{tf.stopgradient( result -in )} &\quad \text{definition of out} \\ \text{\color{red} \textbackslash y}_{(\ll-1)} &= [\text{in}, \text{result}] && \text{inputs to layer}\end{aligned}$$

Then

$$\frac{\partial \text{\\y\\llp}}{\partial \text{\\y}(\ll-1)} = \frac{\partial \text{ out}}{\partial [\text{in,result}]}$$

definitio

$$= \frac{\partial \text{ in} + \text{tf.stopgradient( result -in )}}{\partial [\text{in,result}]}$$

definitio

$$= \frac{\partial \text{ in}}{\partial [\text{in,result}]} + \frac{\partial \text{ tf.stopgradient( result -in )}}{\partial [\text{in,result}]}$$

derivati

$$= [1, 0] + [0, 0]$$

$\frac{\partial \text{ in}}{\partial \text{ in}} = 1$

$$\frac{\partial \text{ in}}{\partial \text{ result}} =$$

$$\frac{\partial \text{sg( .. )}}{\partial ..} =$$

$$[1, 0]$$

Thus

$$\frac{\partial \text{ out}}{\partial [\text{in,result}]} = [1, 0]$$

This above shows

- the derivative flows backwards to `in` undiminished (gradient of `out` w.r.t `in` is 1)
- NO derivative flows backwards to `result` (gradient of `out` w.r.t `result` is 0)

In [2]: `print("Done")`

Done

