

Introduction

TL;DR

- Text is represented as
 - a sequence of integers: an integer index into the list of tokens in the vocabulary
- We want to represent other types of data (e.g., images) as
 - a sequence of integers: an integer index into a "code book": finite list of vectors
- By making the "shape" (sequence) and "type" (integer) compatible between text and other data types
 - We facilitate mixing text and other data types
 - Will enable an implementation of Text to Image as a simple extension of the Language Modeling objective

We now present an Autoencoder with a twist

- the latent representation produced for an input
- is limited to be one member of a *finite list* of vectors
- enabling us to describe the latent by the *integer index* in the list

Why is an integer encoding of an input interesting ?

- It is analogous to the way we treat words (tokens) in Natural Language Processing
 - an index into a finite Vocabulary of words
- This opens the possibility of dealing with sequences that are a *mixture* of text and other data types (e.g., images)

Rather than pre-specifying the finite list, we will *learn* the list by training a Neural Network.

In a subsequent module, we will how

- to use this to create "image tokens" that can be mixed with "language tokens" in a multi-model LLM
- facilitate Text to Image
 - given a description, create an Image matching the description

But there is a significant problem with a Neural Network that learns discrete values

- the network may need to make a "hard" (as compared to "soft") choice
 - a true `if` statement ("hard") versus a "soft" conditional (sigmoid)
 - a Python `dict` ("hard" lookup) versus a "soft" lookup (Context Sensitive Memory)
- "hard" means derivatives are not continuous
- Gradient Descent won't work

There will be a challenge in mapping vectors with continuous-valued elements into token indexes

Vector Quantized Autoencoder

A *Vector Quantized VAE* is a VAE with similarities to PCA.

Given an input $\mathbf{x}$

- vector of continuous values

it creates another continuous vector

- that is an *approximation* of $\mathbf{x}$
- but is chosen from a finite collection of K continuous vectors

It outputs $\mathbf{z}$: the index within the codebook of the approximating vector

The integer index $\mathbf{z}$ can be thought of

- as a token (Categorical variable)
- from a finite Vocabulary of K tokens (the codebook)

That is $\mathbf{z}$

- is a *discrete* representation
- of a continuous vector

For example:

- consider the image below: shape $(h \times w \times 3)$
 - 3 features (integer RGB intensities) over $h * w$ locations
- lots of space when h, w are large

3D image to 2D codes

VQ-VAE

One reason for such a large storage requirement

- pixels are a very surface-level ("syntactic") representation
 - individually: they mean little
 - only when grouped with their neighbors into "concept" ("semantic" representation) do they acquire meaning
 - e.g., line segments, curves, body parts

We can create an alternate representation (middle box of the diagram)

- by processing through a CNN
- D features over $h' * w'$ locations
 - vector of length D at each location
 - "semantic" concepts
 - may span multiple adjacent pixels
- $h \gg h', w \gg w'$

The goal is to

- replace the *continuous vector* at each location
- by one vector e from the learned codebook

$$e \in \mathbf{E}$$

thus obtaining a representation of shape $(h' \times \backslash \mathbf{w}' \times D)$

- which can be reduced to shape $(h' \times \backslash \mathbf{w}')$
- by replacing each vector of length D with its index in $\mathbf{E}$

Details

A VQ-VAE can produce a sequence of integers that encodes many different types of data

- including data with "shape": *non-feature* dimensions
 - image
 - audio

Each "location" in the space of non-feature dimensions

- is a vector consisting only of features

So an image, for example,

- is a collection of feature vectors
- whose locations are arranged into a 2D grid
 - a location is a pair of row number/column number
- the feature vector at each location
 - has 3 features: Red, Green, Blue

We illustrate the VQ-VAE using Image examples.

Here is diagram of a VQ-VAE

- that creates a latent representation of a 3-dimensional image ($w \times h \times 3$)
- as a 2-dimensional matrix of integers
 - each location
 - has an integer index
 - which references a vector of length n in the codebook (with K codes in the codebook)

There is a bit of notation: referring to the diagram should facilitate understanding the notation.

VQ-VAE



The picture: in words

The Encoder

- the *input* image (dog) $\backslash \mathbf{x}$ of shape $(h \times w \times n)$
 - feature vectors of length $n = 3$ at locations defined in a 2D grid of shape $(h \times w)$
- is transformed by a CNN
- into a 3D alternate representation $\mathbf{z}_e(\backslash \mathbf{x})$
 - feature vectors of length D over a 2D grid of shape $S = (h' \times w')$
 - $(h' \times w')$ may differ from $(h \times w)$
 - strided convolution, downsampling

The result is called the *Encoder Output* $\mathbf{z}_e(\backslash \mathbf{x})$

- shape $(S \times D)$
- where each location is a vector of **continuous** values

The feature vector of length D at **each location** of the Encoder Output $z_e(\backslash \mathbf{x})$

- is mapped to one of the possible Embedding Vectors
- defined by the *Codebook* $\mathbf{E}$
 - shape $(K \times D)$
 - a collection of K continuous embedding vectors of length D

This results in an *approximation* $z_q(\backslash \mathbf{x})$ of the Encoder Output $z_e(\backslash \mathbf{x})$

$$z_q(\backslash \mathbf{x}) \approx z_e(\backslash \mathbf{x})$$

- the actual vector at each location in Encoder Output $z_e(\backslash \mathbf{x})$
- is replaced by some vector e_k in $\mathbf{E}$

This approximation is called the *quantized* equivalent of $\backslash \mathbf{x}$, denoted $z_q(\backslash \mathbf{x})$

- shape $(S \times D)$
- where each location is a vector of continuous values
 - one of the K embedding vectors

Since each location in the quantized equivalent $z_q(\mathbf{x})$ is from a finite set $\mathbf{e}$

- we can equivalently represent the information via a spatial arrangement of integers in $[1 : K]$
- the index within $\mathbf{e}$ of the vector at the location

denoted by $q(\mathbf{z}|\mathbf{x})$

- shape $\mathcal{S}$
- where each location is an integer

So mapping the integers in $q(\mathbf{z}|\mathbf{x})$ through the codebook $\mathbf{E}$ results in $z_q(\mathbf{x})$.

The Decoder

- attempts to map the quantized equivalent $z_q(\backslash \mathbf{x})$
- into the reconstructed

$$p(\backslash \mathbf{x} \mid z_q)$$

- that is close to original input $\backslash \mathbf{x}$

$$p(\backslash \mathbf{x} \mid z_q) \approx \backslash \mathbf{x}$$

The picture thus describes an Autoencoder

- where the "bottleneck"
- is caused by quantization rather than dimensionality reduction.

Notation

Notation warning

Our inputs $\backslash \mathbf{x}$ and their alternate representations

- $z_e(\backslash \mathbf{x})$
- $z_q(\backslash \mathbf{x})$
- $q(\backslash \mathbf{z} | \backslash \mathbf{x})$
- $p(\backslash \mathbf{x} | z_q)$

all have non-feature dimensions of shape $S = (n_1 \times n_2 \dots n_{\#S})$.

Rather than explicitly iterating over each location

- we use the unscripted version of the letter
- to refer to the value at a single location

That is, the single location equivalents are denoted

- $\backslash z_e(\backslash x)$
- $\backslash z_q(\backslash x)$
- $q(\backslash z|\backslash x)$
- $\backslash \text{prc} \backslash x \backslash z_q$

Notation summary

term	shape	meaning
S	$(n_1 \times n_2 \dots \times n_{\#S})$	Spatial dimensions of $\#S$ -dimensional input
$\backslash \mathbf{x}$	$\backslash \text{Reals}^{S \times n}$	Input
D		length of latent vectors (Encoder output, Quantized Encoder output, Codebook entry)
$\mathcal{E}$		Encoder function
$\mathbf{z}_e(\backslash \mathbf{x})$	$\mathbb{R}^{S \times D}$	Encoder output over each location of spatial dimension
		$\backslash \mathbf{z}_e(\backslash \mathbf{x}) = \mathcal{E}(\backslash \mathbf{x})$
$\backslash \mathbf{z}_e(\backslash \mathbf{x})$	$\mathbb{R}^D$	Encoder output at a single representative spatial location
		$\backslash \mathbf{z}_e(\backslash \mathbf{x}) = \mathcal{E}(\backslash \mathbf{x})$
K		number of codes
$\mathbf{E}$	$\mathbb{R}^{K \times D}$	Codebook/Embedding
		K codes, each of length D
$e \in \mathbf{E}$	$\mathbb{R}^D$	code/embedding
$\mathbf{z}$	$\{1, \dots, K\}^{S \times D}$	latent representation over each location of spatial dimension
$\backslash \mathbf{z}$	$\{1, \dots, K\}$	Latent representation at a single representative spatial location
		one integer per spatial location
$\$ \backslash \text{qr} \{ \backslash \mathbf{z} \} \backslash \mathbf{x} \}$		integer $\in [1 \dots K]$
		Index k of $e_k \in \mathbf{E}$ that is closest to $\backslash \mathbf{z}_e(\backslash \mathbf{x})$
		$k = \backslash \text{argmin} j \in [1, K] \backslash \mathbf{z}_e(\backslash \mathbf{x}) - \backslash \mathbf{e}_j _2$
		actually: encoded as a OHE vector of length K
$\mathbf{z}_q(\backslash \mathbf{x})$	$\mathbb{R}^{S \times D}$	Quantized $\mathbf{z}_e(\backslash \mathbf{x})$ over each location of spatial dimension
$\backslash \mathbf{z}_q(\backslash \mathbf{x})$	$\mathbb{R}^D$	Quantized $\backslash \mathbf{z}_e(\backslash \mathbf{x})$ at a single representative spatial location
		$\backslash \mathbf{z}_q(\backslash \mathbf{x}) = e_k$ where $\$k = \backslash \text{qr} \{ \backslash \mathbf{z} \} \backslash \mathbf{x} \}$
		i.e, the element of codebook that is closest to $\backslash \mathbf{z}_e(\backslash \mathbf{x})$
		$\backslash \mathbf{z}_q(\backslash \mathbf{x}) \approx \backslash \mathbf{z}_e(\backslash \mathbf{x})$

term	shape	meaning
$\tilde{\mathbf{x}}$	n	Output: reconstructed $\mathbf{x}$
		$\mathcal{P}(\mathbf{x})$
$\mathcal{D}$	$\mathbb{R}^{n'} \rightarrow \mathbb{R}^n$	Decoder

Quantization

Given

- continuous vector $\mathbf{z}_e(\mathbf{x})$ at a single location
- we need to map it to one vector $\mathbf{e}_k$ in codebook $\mathbf{E}$
- creating the approximation vector $\mathbf{z}_q(\mathbf{x})$

The approximation vector $\mathbf{z}_q(\mathbf{x})$ can then be represented as an integer k

$$k = q(\mathbf{z}|\mathbf{x}) \in \{1, \dots, K\}$$

where k is the *index* of a row $\mathbf{e}_k$ in codebook $\mathbf{E}$

$$\mathbf{e}_k = \mathbf{E}_k \in \mathbf{Reals}^D$$

The codebook is also called an *Embedding* table.

k is chosen such that $\mathbf{e}_k$ is the row in $\mathbf{E}$ closest to $\mathbf{z}_e(\mathbf{x})$

$$\begin{aligned} k &= q(\mathbf{z}|\mathbf{x}) \\ &= \operatorname{argmin}_{j \in \{1, \dots, K\}} \|\mathbf{z}_e(\mathbf{x}) - \mathbf{e}_j\|_2 \end{aligned}$$

In our notation:

- $\backslash \mathbf{z}_q(\backslash \mathbf{x})$
- is the quantized version of $\backslash \mathbf{z}_e(\backslash \mathbf{x})$
- where

$$\backslash \mathbf{z}_q(\backslash \mathbf{x}) = \backslash \mathbf{e}_k$$

- such that

$$k = q(\backslash \mathbf{z} | \backslash \mathbf{x})$$

The Decoder tries to invert the codebook entry $\mathbf{e}_k = \mathbf{z}_q(\mathbf{x})$ so that

$$\begin{aligned}\tilde{\mathbf{x}} &= \mathcal{D}(\mathbf{z}_q(\mathbf{x})) \\ &\approx \mathbf{x}\end{aligned}$$

Discussion

Why do we need the CNN Encoder ?

The CNN Encoder takes

- input shape $(h \times w \times 3)$ for RGB images
- output shape $(h' \times w; \times D)$

When $h' = h, w' = w$

- the CNN is *creating* many synthetic features at each spatial location

Eventually, the continuous vector $\mathbf{z}_e(\mathbf{x})$ of length D for one location

- is mapped to an approximate vector of length D from the codebook

Even if we

- limit the range of values of the elements of $\mathbf{z}_e(\mathbf{x})$
- and scale/round them to integers

there are *many* possible vectors of length D

Perhaps we need to expand the feature dimension from 3 to $D \gg 3$

- because a location needs to
 - represent *semantic* information on the image
 - rather than color intensities as for the input
 - in order to concisely encode the many possibilities for vectors of length D

Recall the concept of the *receptive field* for one location in the output of a CNN

- the pixels in the input layer that affect the value at the location
- receptive field size grows with number of CNN layers
 - three layers of a CNN with filter size 3 at each layer is a (7×7) region of the input

So the continuous vector at each location represents *several* input pixels

- and thus is more likely to represent "semantic" concepts of a region
- rather than intensity of a single input location

Although our illustration preserved the non-feature dimensions

- it may be useful for the CNN to *down-sample* spatial dimension S to a smaller S'
- resulting in shorter sequences when we eliminate the non-feature dimensions by "flattening"

For example

- 3 layers of stride 2 CNN layers
- will reduce a 2D image of spatial dimension $(n_1 \times n_2)$
- to spatial dimension $(\frac{n_1}{8} \times \frac{n_2}{8})$

This replaces each $(8 \times 8 \times n)$ *patch* of raw input

- into a single vector of length D
- that summarizes the (8×8) the patch

Why quantize ?

Quantization

- converts the continuous $\mathbf{z}_e(\mathbf{x})$
- into discrete $q(\mathbf{z}|\mathbf{x})$
- representing the approximation $\mathbf{z}_q(\mathbf{x}) \approx \mathbf{z}_e(\mathbf{x})$

The Decoder inverts the approximation.

Why bother when the Quantization/De-Quantization is Lossy ?

One motivation comes from observing what happens if we *quantize and flatten*

- the $\#S'$ -dimensional spatial locations
- to a sequence of integers

Quantizing replaces each patch with a single integer index.

- the integer is the index of an *image token* within a list of K possible tokens

By flattening the quantized higher dimensional matrix of patches, we convert the input

- into a sequence of image tokens
- over a "vocabulary" defined by the codebook $\mathbf{E}$.

This yields an image representation

- similar to the representation of text

Thus, we open the possibility of processing sequences of mixed text and image tokens.

Loss function

The Loss function for the VQ-VAE entails several parts

- Reconstruction loss
 - enforcing constraint that reconstructed image is similar to input
$$\tilde{\mathbf{x}} \approx \mathbf{x}$$
- Vector Quantization (VQ) Loss:
 - enforcing similarity of quantized encoder output and actual encoder output
$$\mathbf{z}_q(\mathbf{x}) \approx \mathbf{z}_e(\mathbf{x})$$
- Commitment Loss
 - a constraint that prevents the Quantization of $\mathbf{z}_e(\mathbf{x})$ from alternating rapidly between codebook entries

The Stop Gradient operator

Some parts of the Loss function will necessitate introduction of a new operator

- *Stop gradient* operator

This operator appears infrequently

- but **importantly**: is the key to incorporating non-differentiable operations into a NN

The issue

- derivatives $\frac{\partial \text{loss}}{\partial \theta}$ are computed under the assumption
 - that "everything else (other than the θ) is held constant"
- but operators that enforce "hard" (vs soft) decisions
 - are discontinuous, hence non-differentiable at the point of discontinuity

In the case of the VAE

- a small change in the parameters of the CNN Encoder
- can cause a big change in the CNN output

$$z_e(\backslash \mathbf{x})$$

- which can cause a change in

$$q(\backslash \mathbf{z} | \backslash \mathbf{x})$$

the *index* in the codebook of some spatial location

- because the index is a *discrete* value, not a continuous value that can change in infinitesimal increments

That is to say:

- Quantization is not differentiable
- It is a "hard" decision: mapping to a discrete value (integer index)

To incorporate the non-differentiable operation

- we use the Stop Gradient operator
- in conjunction with a technique called *Straight Through Estimation*
- that will be explained below

Stop Gradient Operator

On the Forward Pass, Stop Gradient `sg` acts as an Identity operator

$$\text{sg}(\backslash \mathbf{x}) = \backslash \mathbf{x}$$

But on the Backward Pass of Backpropagation: *it stops the gradient* from flowing backwards

$$\frac{\partial \text{sg}(\backslash \mathbf{x})}{\partial \backslash \mathbf{y}} = 0 \text{ for all } \backslash \mathbf{y}$$

Essentially, Stop Gradient "defines" the gradient of a non-differentiable operator to be zero.

So no update is passed backwards through the operation

- thereby ensuring that weights of nodes that are closer to the Input
- don't change
- ensuring "everything else held constant" for those nodes

Reconstruction Loss

The Reconstruction Loss term is our familiar: Maximize Likelihood

- written to minimize the negative of the log likelihood, as usual

$$-\log p(\mathbf{x} | \mathbf{z}_q(\mathbf{x}))$$

It's purpose

- to ensure that the reconstruction is as close to the input as possible

This loss will serve to update parameters for the Encoder and Decoder.

As we will see (section: "Quantization is not differentiable")

- during the backward pass
- the Loss gradient from the quantized $\mathbf{z}_q(\mathbf{x})$
 - which is an input to the Decoder
- flows directly to the (continuous) Encoder output $\mathbf{z}_e(\mathbf{x})$

Effectively, for the purpose of gradient/weight update due to Reconstruction Loss:

$$\mathbf{z}_q(\mathbf{x}) = \mathbf{z}_e(\mathbf{x})$$

- the "quantization" node acts like an Identity function

So the Reconstruction Loss will update all parameters

- **except** for the codebook embeddings

The update of the codebook embeddings is the job of two other loss terms.

Vector Quantization Loss

The Vector Quantization Loss and Commitment Loss are similar.

- differ only in the placement of the Stop Gradient

We need two versions of the same Loss

- in order to satisfy "everything else held constant"
- by holding constant one of
 - Codebook update
 - Encoder weight update
- in each of the cases

Vector Quantization Loss:

$$\| \text{sg}(\mathbf{z}_e(\mathbf{x})) - \mathbf{z}_q(\mathbf{x}) \|$$

where sg is the *Stop Gradient* Operator (details to follow).

Its purpose

- is to ensure that the quantized and un-quantized vectors for an input are **as similar as possible**.

This loss will serve to update the parameters of the codebook $\mathbf{E}$.

- while holding the Encoder output constant
 - by preventing update of its weights: $\text{sg}(\mathbf{z}_e(\mathbf{x}))$

Commitment Loss

Commitment Loss:

$$\|\mathbf{z}_e(\mathbf{x}) - \text{sg}(\mathbf{z}_q(\mathbf{x}))\|$$

It is similar to the Vector Quantization loss except for the placement of the Stop Gradient operator.

Its purpose

- is to ensure that the quantized and un-quantized vectors for an input are **as similar as possible**.

This loss serves to update the weights of the Encoder

- while holding the Codebook constant: $\text{sg}(\mathbf{z}_q(\mathbf{x}))$
- by preventing updates to **E**

Total Loss

Loss function

$$\begin{aligned} \text{loss}(\mathbf{x}, \mathcal{D}(\mathbf{e})) = & \|\mathbf{x} - \mathcal{D}(\mathbf{e})\|_2^2 && \text{Reconstruction Loss} \\ & + \|\text{sg}[\mathcal{E}(\mathbf{x})] - \mathbf{e}\|_2^2 && \text{VQ loss, codebook loss: train codebook} \\ & + \beta \|\text{sg}[\mathbf{e}] - \mathcal{E}(\mathbf{x})\|_2^2 && \text{Commitment Loss: force } \mathcal{E}(\mathbf{x}) \text{ to} \\ & \text{where } \mathbf{e} = \mathbf{z}_q(\mathbf{x}) \end{aligned}$$

Need the stop gradient operator sg to control the mutual dependence

- of the Encoder output $\mathcal{E}(\mathbf{x})$ and the chosen code $\mathbf{e}$
-

Quantization is not differentiable

There is a subtle but important problem.

The Quantization operation

$$\begin{aligned}k &= q(\mathbf{z}|\mathbf{x}) \\ &= \operatorname{argmin}_{j \in \{1, \dots, K\}} \|\mathbf{z}_e(\mathbf{x}) - \mathbf{e}_j\|_2\end{aligned}$$

is not differentiable because of argmin

`argmin` is a problematic operation because

- it contains a "hard choice" so is not differentiable
 - output may change dis-continuously from index k to index $k' \neq k$
 - for small changes in the input
 - not continuous as the point of change
- it may also be non-deterministic
 - when minimum value occurs at more than one index
 - when $\backslash e_k = \backslash e_{k'}$ for $k \neq k'$

There is a work-around

- implement a VectorQuantizer layer
- using a *Straight Through Estimator*

The Straight Through Estimator. when applied to a node

- uses the Stop Gradient operator
- in conjunction with a specific construction
- to cause the Loss Gradient to *pass through the node unchanged* during back propagation
 - essentially treating the operation of the node as the Identity operation
 - instead of its true computation

We will learn the details in the module [Straight Through Estimator](#)
[\(Straight_Through_Estimator.ipynb\)](#).

We see this in the [Colab \(https://keras.io/examples/generative/vq_vae/\)](https://keras.io/examples/generative/vq_vae/) implementation of Vector Quantization (the VectorQuantizer layer)

```
class VectorQuantizer(layers.Layer):  
    ...  
    def call(self, x):  
    ...  
        # Straight-through estimator.  
        quantized = x + tf.stop_gradient(quantized - x)
```

Code similar to the [VectorQuantizer of the paper's authors \(https://github.com/deepmind/sonnet/blob/v1/sonnet/python/modules/nets/vqvae.py\)](https://github.com/deepmind/sonnet/blob/v1/sonnet/python/modules/nets/vqvae.py)

Code

Here (https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/vq_vae.ipynb#scrollTo=LWYJf1MYvzap) is a Colab notebook.

Creating synthetic examples

Consider an example that is an Image.

It is a structured arrangement of feature vectors

- a 2D grid
- each element of the grid is a vector of features

In general: our examples may have an arbitrary number of dimensions.

By convention, we will refer

- to the last dimensions as the "feature" dimension
- all the preceding dimensions as *non-feature* dimensions
 - for an Image: *spatial* dimensions

Denoting $\#S$ as the number of non-feature dimensions and $n_{(0)}$ as the number of features, examples have shape

$$(n_1 \times \dots \times n_{\#S} \times n_{(0)})$$

where $\#S$ denotes the number of non-feature dimensions ($\#S = 2$ for an Image)

The Encoder creates an output $\mathbf{z}_e(\mathbf{x})$ of shape
 $(n_1 \times \dots \times n_{\#S} \times n_e)$
which is quantized into $\mathbf{z}_q(\mathbf{x})$ of shape
 $(n_1 \times \dots \times n_{\#S} \times 1)$

That is: the feature dimension of $\mathbf{z}_q(\mathbf{x})$ is a single integer index into a learned codebook.

In order for us to generate synthetic examples

- we must create a Tensor $\mathbf{z}_q(\mathbf{x})$ of $(n_1 \times \dots \times n_{\#S} \times 1)$ integers (*latents*)
- feed this to the Decoder
- get a synthetic example as output

Learning the distribution of latents

But we can't create $\mathbf{z}_q(\mathbf{x})$ *completely at random*.

- The elements at adjacent locations in the non-feature dimensions may not be **independent**

For example: consider an Image

- the pixels in an image are related to one another

We must learn a distribution of $\mathbf{z}_q(\mathbf{x})$ that respects the dependencies.

One solution

- flatten $\backslash \mathbf{z}_q(\backslash \mathbf{x})$
- from shape $(n_1 \times \dots \times n_{\#S} \times 1)$
- into a sequence
 $\backslash \mathbf{z}_{(1)}, \backslash \mathbf{z}_{(2)}, \dots, \backslash \mathbf{z}_{(n_1 * n_2 \dots * n_{\#S})}$
of integer indices.

We can then learn the distribution of the latents

- through auto-regressive modeling of the sequence $\mathbf{z}$
 $\Pr(\mathbf{z}_{(k+1)} | \mathbf{z}_{(1)}, \dots, \mathbf{z}_{(k)})$

This is just like the Language Model objection for NLP.

Aside

Learning a distribution is less restrictive than *assuming* a distribution

For the case of "structured" examples: we have no choice.

- adjacent elements are not independent

But recall that, for the VAE, we *assumed* that the latents came from a Normal distribution.

Use Auto-regressive modeling is a nice "trick" to learn distributions rather than having to assume a "convenient" functional form.

PCA versus VQ-VAE

The common element in the design of any Autoencoder method is

- to create a latent representation $\mathbf{z}$ of input $\mathbf{x}$
- such that $\mathbf{z}$ can be (approximately) inverted to reconstruct $\mathbf{x}$.

Principal Components Analysis is a type of Autoencoder that produces a latent representation $\mathbf{z}$ of $\mathbf{x}$

- $\mathbf{x}$ is a vector of length n : $\mathbf{x} \in \mathbb{R}^n$
- $\mathbf{z}$ is a vector of length $n' \leq n$: $\mathbf{z} \in \mathbb{R}^{n'}$

Usually $n' \ll n$: achieving *dimensionality reduction*

This is accomplished by decomposing $\mathbf{x}$ into a weighted product of n *Principal Components*

- $\mathbf{V} \in \text{Reals}^{n \times n}$
 - $\mathbf{x} = \mathbf{z}' \mathbf{V}^T$
 - where $\mathbf{z}' \in \text{Reals}^n$
 - rows of $\mathbf{V}^T$ are the components

So $\mathbf{x}$ can be decomposed into the weighted sum (with $\mathbf{z}'$ specifying the weights)

- of n component vectors
- each of length n

That is:

- $\mathbf{V}^T$ is a basis space that can define all vectors in n -dimensional space
- with loadings $\mathbf{z}'$ on the basis vectors.

To illustrate

- think of the familiar basis space $I \in \text{\textcolor{red}{Reals}}^{n \times n} = \text{diagonal}(n)$
 - $\mathbf{x}$ are the "coordinates" of a point within the n -dimensional space spanned by I
- $\mathbf{z}'$ are the "coordinates" of *the same point* but described in basis space $\mathbf{V}^T$

Since $\mathbf{z}' \in \text{\textcolor{red}{Reals}}^n$: there is **no** dimensionality reduction just yet.

One can also view $\mathbf{V}^T$ as a kind of *codebook*

- any $\mathbf{x}$ can be represented (as a linear combination) of the *codes* (components) in $\mathbf{V}^T$

$$\mathbf{x} = \mathbf{z}' \mathbf{V}^T$$

$\mathbf{z}'$ is like a translation of $\mathbf{x}$, using $\mathbf{V}$ as the vocabulary.

- weights in the codebook
- rather than weights in the standard basis space $I \in \text{Reals}^{n \times n} = \text{diagonal}(n)$

$$\mathbf{x} = \mathbf{x} I$$

So can we take the PCA of images ?

Yes !

We did this for the MNIST digits and Fashion MNIST

- each basis vector is a pseudo digit/clothing item
- which can be combined
- into an approximation of any digit/clothing item in the *training* dataset

But these were small and simple examples (28×28 pixels)

We can do this for more complex images


- Faces

See Eigenfaces (https://scikit-learn.org/stable/auto_examples/applications/plot_face_recognition.html).

eigenface 0



eigenface 1



eigenface 2



eigenface 3



Thus, the "eigen faces" are a "codebook" of basic faces.

So a face in the training dataset can be encoded

- as a linear combination of $r \leq n$ codebook entries

Comparing PCA to VQ-VAE

- PCA chooses $r \leq n$ items from the codebook
 - versus $r = 1$ for VQ-VAE
- The weights $\mathbf{z}'$ (vector of length r) are *continuous*
 - versus a single integer index for the $\mathbf{z}$ of VQ-VAE

Also

- PCA is often used for dimensionality reduction
 - so the size of the codebook $r < n$ is relatively small
- VQ-VAE codebooks are larger
 - a couple hundred to a couple thousand entries
 - codebook entries are typically of length 128 to 512

So PCA is an alternative to VQ-VAE for representing images

- for **some** use cases
- but **not** for the use case of creating a discrete representation of the image

In [2]: `print("Done")`

Done

